Predicting the Next Category and Region of a Visit

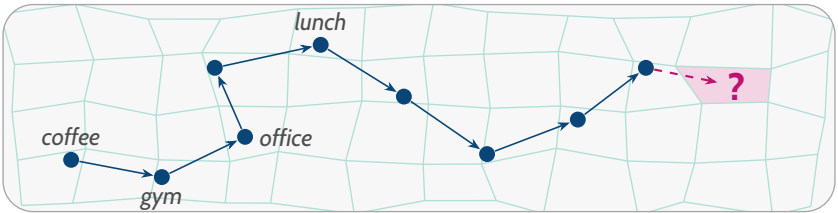
A Check-in-Level Multi-Task Study on Mobility Data

Vitor Hugo de Oliveira Silva Germano Barcelos dos Santos
Fabrício Aguiar Silva

NESPeD-LAB · Universidade Federal de Viçosa, Brazil

MobiWac 2026 · Paris · October 26, 2026

One person's day, and two questions



What kind of place?

Next category: one of 7 types, such as food, shopping, nightlife. → *what the user will want*

Which neighborhood?

Next region: its census tract (Istanbul: *mahalle*), 520 to 8,501. → *where to prepare*

Check-in: a visit shared on a location-based social network such as Gowalla: a place (a *point of interest*, POI) and a time (Cho 2011). Uses: caching, capacity, demand anticipation (Bastug 2014; Moura 2025). A tract is not a radio cell: we claim nothing about cells or handover.

Two models, or one?

- **Two dedicated models**: one per question; two models to train, version and serve.
- **One joint model**: one artifact to deploy, and the tasks might help each other. But shared parameters can settle on a compromise good for neither task (Caruana 1997).

Multi-task models in this field use category and region as auxiliary signals for next-place prediction (MCMG, Sun 2024; HMT-GRN, Lim 2022). To our knowledge, fine-grained region as an end target of equal standing is underexplored.

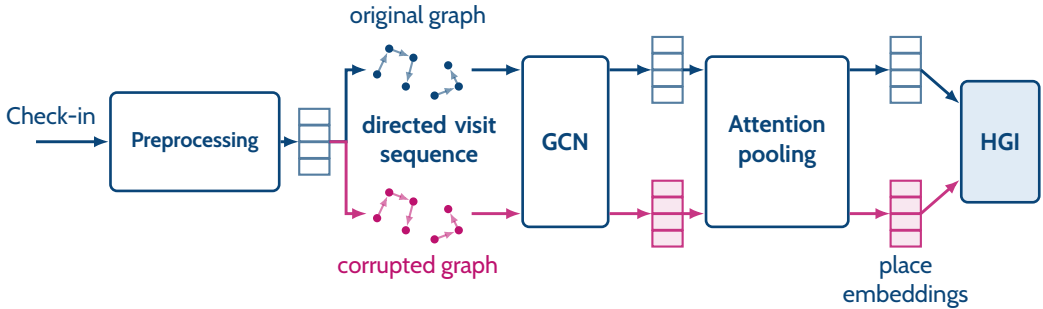
The question

Can one model predict both the category and the region of a user's next visit, and what does it cost?

Short answer: yes, at a measured price.

Check2HGI: a fourth level below the place

Place embeddings: one vector per place, DGI (Veličković 2019), HGI (Huang 2023). Two visits to one café look identical.



- **Each visit carries** its place's category, the hour and day of week, and four elapsed-time values → one **64-d** vector per visit, **the vector our model reads**; pooled per place, it feeds the levels above.
- Trained **without the task labels**, and **built from the past alone**: edges between consecutive visits run forward in time only.

How it was measured

Regions	census tracts (Istanbul: *mahalle*), **520** to **8,501**
Split	by user: the prediction models never train on a test user's visits
Scored	category: **macro-F1** (always-majority: 5.7 to 7.3)
	region: **Acc@10**, is the true region in the top ten? (random: at most about 2%)
Repeated	**20 trained models** per configuration (4 seeds $\times$ 5 folds)
Decided	written down before any result was seen:
	category: is the joint model **better**? (superiority)
	region: is it **not worse** by more than **2 Acc@10 points**? (non-inferiority, TOST) (Lakens 2017)

Result 1: the representation, at every dataset

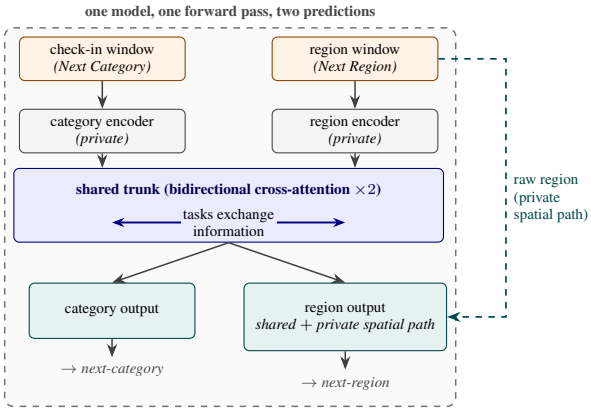
Next category, single-task model, seed 0 — only the input changes

Dataset	Check-in level	Place level (HGI)	Δ
Alabama	**30.77** $\pm$1.16	29.15 $\pm$0.84	$+1.62$
Arizona	**34.51** $\pm$1.15	31.93 $\pm$1.08	$+2.58$
Istanbul	**35.35** $\pm$0.89	32.54 $\pm$1.03	$+2.81$
Florida	**37.36** $\pm$0.42	37.13 $\pm$0.41	$+0.23$
California	**35.62** $\pm$0.47	34.74 $\pm$0.41	$+0.88$
Texas	**36.32** $\pm$0.51	35.33 $\pm$0.48	$+0.99$

macro-F1; five matched folds, $\pm$ is the fold sd.

Larger at the three smaller datasets ($+1.6$ to $+2.8$), under one point at the three largest. A paired test separates the two at five of six; not at Florida.

One model for two tasks



Multi-task learning: one model does several tasks at once by sharing most of its parts.

The risk: what the tasks share can settle on a compromise worse for each.

Here: each task has its own input and encoder. In the shared trunk, attention lets each task read the other's features while keeping its own weights: **they share by exchanging information**.

A private spatial path keeps what is purely spatial out of the exchange.

One model, two answers: the price

	Next region (Acc@10)			**Next category** (macro-F1)		
	Ded.	Joint	Δ	Ded.	Joint	Δ
Alabama	**70.12**	69.24	−0.87	**30.77**	30.59	−0.19
Arizona	**59.48**	59.04	−0.44	**34.57**	**34.57**	−0.00
Istanbul	**75.16**	75.08	−0.08	35.34	**35.42**	+0.08
Florida	**76.69**	76.54	−0.16	37.35	**37.55**	+0.19
California	63.48	**64.54**	+1.06	**35.63**	**35.63**	−0.00
Texas	64.94	**66.15**	+1.21	**36.33**	36.19	−0.13

Ded.: two dedicated models; Joint: one model. **Green**: outperforms (Holm-corrected). Bold: the larger. **Region**: the other four are deficits, not ties, inside the 2-point margin. **Category**: every interval within half a point; only Florida survives correction. Δ: paired mean over four seeds (can differ from the rounded columns).

The trade is a measured one and not a free substitution.

And against published systems

Next category (macro-F1)

	POI-RGNN	Ded.	Joint
AL	23.80	**30.77**	30.59
AZ	27.64	**34.57**	**34.57**
IST	30.12	35.34	**35.42**
FL	33.35	37.35	**37.55**
CA	30.71	**35.63**	**35.63**
TX	32.08	**36.33**	36.19

Next region (Acc@10)

	HMT-GRN	ReHDM	STAN	Ded.	Joint
AL	63.19	65.38	60.72	**70.12**	69.24
AZ	52.11	53.00	49.86	**59.48**	59.04
IST	68.06	69.33	61.86	**75.16**	75.08
FL	72.76	64.49	72.99	**76.69**	76.54
CA	59.59	50.26	58.61	63.48	**64.54**
TX	61.39	48.81	61.75	64.94	**66.15**

Both our models are above every system we compared against; the joint model by at least **4.11** (category) and **3.55** (region). That margin includes the representation's advantage, so it is not evidence for sharing. Bold: best in row; underline: second.

All four run on our check-in data; only HMT-GRN on our windows and folds. HMT-GRN trimmed to its region target; STAN's output adapted to rank regions. POI-RGNN (Capanema 2023), HMT-GRN (Lim 2022), STAN (Luo 2021), ReHDM (Li 2025). Details: backup B7.

What a mobility-aware service would receive

The category says **what** the user will want; the region says **where** to prepare, before the user arrives: e.g., caching content or provisioning capacity by neighborhood (Bastug 2014).

	AL	AZ	IST	FL	CA	TX
regions	1,109	1,547	520	4,703	8,501	6,553
Transition table, no learning	62.26	51.23	65.06	72.47	59.09	60.10
Joint model	69.24	59.04	75.08	76.54	64.54	66.15
Gain	+6.98	+7.81	+10.02	+4.07	+5.45	+6.05

Acc@10: is the true next region among the ten most likely? Floor: first-order Markov over region transitions, same windows, user-grouped folds, a different partition.

Time to act: half of the next visits come within 0.4 h (Florida) to 5.5 h (Istanbul); 5 to 27% come more than 3 days later.

We built and evaluated no such service: this is motivation, not a measured service result.

What it costs, and three limits

THE COST

Parameters, joint vs **the two dedicated models combined**:

AL 4.2 M vs 1.9 M
CA 5.2 M vs 2.8 M

More compute per pass.

Gain: one model to train, version and deploy.

THREE LIMITS

1. The representation **saw all users once**. Rebuilt per fold from training users only, the dedicated models move by at most 0.17 points (one seed, the datasets checked).
2. The epoch is chosen on the fold that is scored: **absolute scores are optimistic**; joint vs dedicated much less.
3. The joint model is **bigger**: its region gains may come from size, not sharing. The one-model result stands either way.

Conclusion

Can one model predict both, and what does it cost?

Yes. Region: outperforms the dedicated model at Texas and California; at most 0.87 behind elsewhere, inside the 2-point margin. **Category**: within half a point. Above every published system we compared against. **Price**: a larger model. **Gain**: one artifact to deploy.

- **Built**: a check-in-level representation that extends hierarchical graph infomax from the place to the individual visit, a fixed per-visit vector other models can reuse; and one model for both tasks.
- **Next**: the **exact next place** as a third target on the same representation; an inductive variant for unseen places and users; denser data and more cities; how far a miss lands, inside a service.

What a user will do next, and where, from one model instead of two.

Thank you

Questions?

Code (model, representation, baselines, statistical tests)
`github.com/VitorHugoOli/PoiMtlNet/tree/mobiwac`

Data (public) · category-annotated Gowalla · Massive-STEPS (Istanbul)

Supported by **FAPEMIG**, **CNPq** and **CAPES**.

`vitor.h.oliveira@ufv.br`

B1 · Why census tracts

- Predicting over a partition of the map is standard, usually with a **grid cell**; we use **official neighborhood-scale units**;
- Next region is classification over all the dataset's regions:

	Istanbul	AL	AZ	FL	TX	CA
Regions	520	1,109	1,547	4,703	6,553	8,501

B2 · The margin, and the tests

- **Registered before any result was read**: superiority for category, non-inferiority for region with a margin of **2 Acc@10 points**, paired Wilcoxon;
- **Why 2**: a service that acts on which neighborhood will be busy would not behave differently below it; the seed-to-seed sd of the paired difference is 0.02 to 0.16;
- **Reported**: paired t on the four per-seed means, 90% interval, Holm across the six datasets per task; at four seeds the exact Wilcoxon cannot go below 0.0625. Both tests are in the code release;
- Category has **no registered equivalence margin**: a difference that fails superiority is unresolved, never a match;
- The two region gains are **secondary results**, outside the registered plan.

B3 · The Markov floor

- A first-order Markov chain over region transitions, no learning: **51.23 to 72.47** Acc@10;
- The joint model exceeds it by $+4.1$ **to** $+10.0$ at every dataset; the narrowest is **Florida,** $+4.07$.

B4 · Can a visit vector see its target?

- Each visit node draws on the **visits that precede it**: the vector encodes the visit with its own history;
- The graph does **not pass information from a later visit back to an earlier one**, in training or at readout;
- That is what keeps a node from carrying a feature of the target it is used to predict;
- The graph is trained without the task labels: it never sees the next category or the next region.

B5 · One model instead of two

- **4.2 M vs 1.9 M** parameters at Alabama, **5.2 M vs 2.8 M** at California, against **the two dedicated models combined**;
- One pass costs more compute than running the two small models;
- The gain is operational: one artifact to train, version and deploy; one pass returns both answers.

B6 · Then what is multi-task learning for?

- On category, the representation moves the score **more than the choice between one model and two**;
- Whether the exchange between the tasks adds anything beyond the representation is not separated by the evidence here;
- What one model gives for certain is operational: one artifact, two answers.

B7 · The external systems

- **HMT-GRN** (Lim 2022) — primary region comparison: our windows and folds, one seed, best validation epoch; its multi-task structure kept, the graph components and hierarchical beam search removed, a per-fold region-transition prior added (our own reported models do not use it); not the complete published system;
- **STAN** (Luo 2021) — re-implemented with its own representations and splits from the raw check-ins, output adapted to rank regions, seed 0;
- **ReHDM** (Li 2025) — its own published chronological split; single seed at TX and CA;
- **POI-RGNN** (Capanema 2023) — next category, re-implemented, non-overlapping windows.

B8 · More regions, more gain?

	Regions	Region Δ (Acc@10)
Texas	6,553	$+1.21$
California	8,501	$+1.06$

California has more regions than Texas and a **smaller** gain.

An observation, not a law.

B9 · Is it just the raw features?

Control: the place vector + the same raw per-visit features (category one-hot, hour, day of week), same single-task model. Run at three datasets.

	Control gain	Check-in-level gain
Alabama	$+1.73$	$+1.62$
Arizona	$+1.70$	$+2.58$
Florida	$+1.02$	$+0.23$

macro-F1 over the place-level representation, category axis.

"Most of the category difference is therefore already available in the raw per-visit features, and this control does not separate what the check-in-level representation adds beyond them on this axis."

B10 · The intervals

Joint minus dedicated, 90% interval of the paired difference

	Next region (Acc@10)		**Next category** (macro-F1)	
	Δ	90% CI	Δ	90% CI
Alabama	-0.87	$[-1.00, -0.75]$	-0.19	$[-0.33, -0.04]$
Arizona	-0.44	$[-0.62, -0.25]$	-0.00	$[-0.04, +0.03]$
Istanbul	-0.08	$[-0.16, -0.002]$	$+0.08$	$[+0.01, +0.15]$
Florida	-0.16	$[-0.19, -0.13]$	$+0.19$	$[+0.14, +0.25]$
California	$+1.06$	$[+1.03, +1.08]$	-0.00	$[-0.03, +0.02]$
Texas	$+1.21$	$[+1.13, +1.29]$	-0.13	$[-0.19, -0.08]$

Green: outperforms, Holm-corrected — region at Texas ($p = 0.00013$) and California ($p < 10^{-4}$), 20 of 20 folds; category at Florida ($p = 0.011$), 19 of 20.

Every region interval of the other four lies **entirely below zero**: deficits, not ties (Istanbul only just).

B11 · References

Bastug, Bennis, Debbah. Living on the Edge: The Role of Proactive Caching in 5G Wireless Networks. *IEEE Commun. Mag.*, 2014.

Capanema et al. Combining recurrent and Graph Neural Networks to predict the next place's category. *Ad Hoc Netw.*, 2023.

Caruana. Multitask Learning. *Mach. Learn.*, 1997.

Cho, Myers, Leskovec. Friendship and Mobility: User Movement in Location-Based Social Networks. *KDD*, 2011.

Huang et al. Learning Urban Region Representations with POIs and Hierarchical Graph Infomax. *ISPRS J. Photogramm. Remote Sens.*, 2023.

Lakens. Equivalence tests: A practical primer for t tests, correlations, and meta-analyses. *Soc. Psychol. Personal. Sci.*, 2017.

Li et al. Beyond Individual and Point: Next POI Recommendation via Region-aware Dynamic Hypergraph with Dual-level Modeling. *IJCAI*, 2025.

Lim et al. Hierarchical Multi-Task Graph Recurrent Network for Next POI Recommendation. *SIGIR*, 2022.

Luo, Liu, Liu. STAN: Spatio-Temporal Attention Network for Next Location Recommendation. *WWW*, 2021.

Moura, Aquino, Loureiro. On the Design of Mobility-Aware Systems: A Tourist's Perspective. *MSWiM*, 2025.

Sun et al. A Multi-channel Next POI Recommendation Framework with Multi-granularity Check-in Signals (MCMG). *ACM TOIS*, 2024.

Veličković et al. Deep Graph Infomax. *ICLR*, 2019.

Vielhaus et al. Handover Predictions as an Enabler for Anticipatory Service Adaptations in Next-Generation Cellular Networks. *MobiWac*, 2022.

B12 · The datasets

Dataset	Source	Check-ins	Users	POIs	Regions	Windows	Max len	Avg len	Dens. (%)	Maj. (%)
Istanbul	Massive-STEPS	462,615	23,694	29,816	520	271,666	817	19.5	0.065	33.4
AL	Gowalla	113,846	3,858	11,848	1,109	96,326	3,835	29.5	0.249	34.2
AZ	Gowalla	236,450	7,869	20,666	1,547	200,895	5,589	30.1	0.145	34.0
FL	Gowalla	1,407,034	21,052	76,544	4,703	1,274,418	16,679	66.8	0.087	24.7
TX	Gowalla	4,089,892	38,644	160,938	6,553	3,830,414	42,300	105.8	0.066	31.0
CA	Gowalla	3,171,380	37,090	169,145	8,501	2,925,466	14,855	85.5	0.051	32.7

Lengths: check-ins per user. Density: $100\times$ check-ins / (users $\times$ POIs). Majority: share of next-visit labels in the most common category. Check-ins, users and POIs count the full corpus; windows are counted after the ten-visit user filter.